

Venkata Anudeep Adiraju

Data Scientist — Feature Engineering — ETL Pipelines — Statistical & Machine Learning
Austin, TX | (917) 291-7024 | anudeep.adiraju@gmail.com | LinkedIn | GitHub | Portfolio

PROFESSIONAL SUMMARY

Data Scientist with healthcare analytics expertise, specializing in **feasibility studies**, **feature engineering**, and **ETL pipeline** development for clinical and operational use cases. Built production systems processing 5M+ health records using **Python, SQL, and Java**, deployed across **AWS, Azure, and GCP** cloud platforms. Expert in applying **statistical and machine learning techniques** to high-dimensional healthcare data, creating **Grafana dashboards** for business stakeholders, and deriving predictive models that drive clinical and financial outcomes.

TECHNICAL SKILLS

Programming & Data Science: Python (pandas, NumPy, scikit-learn), Java, SQL, R, Statistical Analysis, Predictive Modeling

Machine Learning & Statistics: Statistical Techniques, Model Evaluation, Model Selection, Feature Engineering, A/B Testing, PyTorch

Data Engineering: ETL Pipelines, Data Transformation, Feature Extraction, Data Quality Frameworks, Apache Spark, Batch Processing

Cloud & Infrastructure: AWS (EC2, S3, Lambda, Redshift), Microsoft Azure, Google Cloud (GCP), OCI Cloud, Docker, Kubernetes

Dashboards & Visualization: Grafana, OpenSearch, Oracle Analytics Cloud, Tableau, Plotly, Streamlit, Business Intelligence Reporting

AI/ML Tools: LLMs, RAG Systems, BERT, Vector Databases, MLflow, Model Serving (FastAPI), NLP

EDUCATION

M.S. Computer Engineering | University of Texas at San Antonio | GPA: 3.70/4.0

Dec 2025

Relevant Coursework: Machine Learning, Statistical Methods, Database Systems, Cloud Computing, Advanced Algorithms

B.Tech Computer Science | HITAM, India | GPA: 3.8/4.0

Sep 2023

DATA SCIENCE EXPERIENCE

Data Scientist – Healthcare Analytics | University of Texas at San Antonio

Jan 2024 – Present

- Conducted **feasibility studies** to assess data availability and quality for clinical research; designed **feature extraction pipelines** in **Python and SQL** processing 5M+ health records, identifying modeling requirements that facilitated 3 new predictive use cases for business stakeholders
- Designed **ETL pipelines** on **AWS** integrating heterogeneous healthcare data sources (clinical trials, patient records, claims data); built data transformation workflows achieving **42%** query performance improvement supporting downstream statistical analysis
- Applied **statistical and machine learning techniques** to build explanatory features from 55K+ healthcare conversation pairs; performed iterative **model selection and evaluation** achieving **82%** prediction accuracy on clinically important treatment outcomes
- Built **Grafana dashboards** and **OpenSearch** visualizations for business teams and clinical leaders; created real-time monitoring tracking system health, **ETL pipeline** latency, and model performance metrics across **500K+** daily requests
- Developed **Java**-based data ingestion service extracting features from unstructured clinical documents; implemented NLP preprocessing pipeline on **Google Cloud** reducing manual data processing time by **91%** for research teams
- Collaborated with **business stakeholders** to identify data analysis requirements; proposed new uses for existing datasets, delivering predictive models that reduced patient dropout rates by **23%** through early intervention recommendations
- Supported iterative **model experimentation** workflows using MLflow on **Azure**; evaluated 12+ candidate models using cross-validation and statistical significance testing, selecting best fit algorithms for financially and operationally important healthcare use cases

Data Science Intern | Aimlytics Technologies

Mar 2022 – May 2022

- Built **ETL pipelines** in **Python and SQL** on **Azure** cloud processing behavioral data for **100K+** daily users; created features from user interaction logs supporting predictive churn modeling with **85%** accuracy
- Applied **statistical and machine learning techniques** to build features from high-dimensional user data; delivered insights to **business teams** driving **40%** improvement in API response times through data-driven optimization
- Created **analytics dashboards** tracking A/B test results and user behavior metrics; automated reporting pipelines delivering daily insights to stakeholders enabling data-driven product decisions

DATA SCIENCE PROJECTS

QuitTxt Clinical Data Science Platform | Healthcare Feature Engineering | GitHub

- Led **feasibility study** for multi-institutional smoking cessation research; built **ETL pipelines** in **Python and SQL** processing clinical trial data from UTSA and UT Health San Antonio with HIPAA-compliant data governance
- Designed features from unstructured health data using **statistical and machine learning techniques** on **AWS**; built RAG pipeline with vector database processing **10K+** medical documents for evidence-based **predictive modeling**
- Built **Grafana monitoring dashboard** for clinical research teams tracking patient engagement metrics; enabled **business stakeholders** to identify at-risk participants, improving study retention by **18%**

SCRAPE Neural Feature Extraction Engine | Scalable ETL Pipeline | Live Demo

- Architected scalable **ETL pipeline** in **Python** for **feature extraction** from unstructured web data; processed **1K+** URLs/hour using distributed task queue on **GCP** with intelligent rate limiting and error handling
- Built content **feature engineering** workflow with BM25 filtering; achieved **95%** extraction accuracy through iterative **model evaluation**, demonstrating ability to industrialize collection from heterogeneous data sources

CERTIFICATIONS & ACHIEVEMENTS

Cloud Platforms: AWS Certified Cloud Practitioner, Azure Fundamentals (AZ-900), Azure AI Fundamentals (AI-900)

AI/ML: TensorFlow Developer Certificate | **Recognition:** DGX Spark Hackathon Winner – GPU-accelerated ML pipelines